

# Maintaining Trace Context in Hybrid Async/Non-Async Rust Applications

This document addresses the challenge of maintaining a single, continuous trace across an event-driven Rust application that mixes asynchronous (Axum, WebSockets) and multi-threaded non-asynchronous processing, linked by unbounded channels.

## 1. The Challenge: Bridging Context Gaps

Your architecture involves several transitions where trace context can be lost if not explicitly handled:

1. **Async to Non-Async (WebSocket Server -> Processing Queue):** The initial event comes from an async Axum handler, but its processing starts on a non-async worker thread.
2. **Non-Async Thread Hops (Processing Thread A -> Processing Thread B/C):** Events are dispatched between different non-async worker threads.
3. **Non-Async to Async (Processing -> External WebSocket Client):** Processed data is sent from a non-async worker to an async WebSocket client for external communication.

The core solution lies in **explicitly propagating the OpenTelemetry Context** through your channel messages.

## 2. Key Concepts for Context Propagation

- **opentelemetry::Context:** This is the central object that carries trace information (like trace ID, span ID of the parent, baggage) across different parts of your application.
- **Span::current().context():** When a tracing span is active, you can retrieve the underlying OpenTelemetry Context associated with it. This context contains the information about the currently active span.
- **Context::attach() / Context::with\_current():** These methods allow you to make a specific `opentelemetry::Context` the "current" context for the executing thread. Any new tracing spans created while this context is active will automatically become children of the span represented by that context.
- **TextMapPropagator:** While `opentelemetry::Context` is an in-memory object, you often need to serialize it to pass it across boundaries (like network requests or, in your case, channels). The `TextMapPropagator` (specifically `TraceContextPropagator` for W3C Trace Context) helps inject context into and extract it from a `HashMap<String, String>` (or similar text-based carrier).

## 3. Implementation Strategy

The strategy involves modifying your event structures to carry the trace context and then

explicitly injecting and extracting this context at each channel boundary.

### Step 3.1: Modify Your Event Message Structure

Your event messages (the ones passed through the unbounded channels) need a new field to carry the trace context. A `HashMap<String, String>` is a good choice as it's easily serializable and compatible with `TextMapPropagator`.

// In your common types or event definitions file (e.g., `src/types.rs`)

```
use std::collections::HashMap;
use opentelemetry::Context;
use opentelemetry::propagation::{TextMapPropagator, Extractor, Injector};
use opentelemetry::sdk::propagation::TraceContextPropagator; // W3C Trace Context
```

```
/// A custom struct to hold your event data and the trace context.
#[derive(Debug, Clone)] // Clone is often useful for channel messages
pub struct MyEvent {
    pub id: String,
    pub payload: String,
    /// This HashMap will carry the serialized OpenTelemetry trace context.
    pub trace_context_carrier: HashMap<String, String>,
}
```

```
/// Helper struct for injecting/extracting context into/from a HashMap.
/// Implements `Injector` and `Extractor` traits from OpenTelemetry.
pub struct HashMapCarrier<'a>(&mut HashMap<String, String>);
```

```
impl<'a> Injector for HashMapCarrier<'a> {
    fn set(&mut self, key: &str, value: String) {
        self.0.insert(key.to_string(), value);
    }
}
```

```
impl<'a> Extractor for HashMapCarrier<'a> {
    fn get(&self, key: &str) -> Option<&str> {
        self.0.get(key).map(|s| s.as_str())
    }

    fn keys(&self) -> Vec<&str> {
        self.0.keys().map(|s| s.as_str()).collect()
    }
}
```

```
/// Injects the current OpenTelemetry Context into a new HashMap.
```

```

pub fn inject_current_context() -> HashMap<String, String> {
    let propagator = TraceContextPropagator::new();
    let parent_context = Context::current(); // Get the current active context
    let mut carrier = HashMap::new();
    propagator.inject_context(&parent_context, &mut HashMapCarrier(&mut carrier));
    carrier
}

/// Extracts an OpenTelemetry Context from a HashMap.
pub fn extract_context_from_carrier(carrier: &HashMap<String, String>) -> Context {
    let propagator = TraceContextPropagator::new();
    propagator.extract_context(&HashMapCarrier(carrier))
}

```

### Step 3.2: Initialize OpenTelemetry and tracing Subscriber (Main)

Ensure your main.rs (or application entry point) has the OpenTelemetry and tracing setup as discussed previously. This provides the global tracer and subscriber.

#### Cargo.toml dependencies (reiterated for clarity):

```

[dependencies]
tracing = "0.1"
tracing-subscriber = { version = "0.3", features = ["env-filter", "fmt"] }
tracing-opentelemetry = "0.23"
opentelemetry = { version = "0.22", features = ["trace", "metrics"] }
opentelemetry-sdk = { version = "0.22", features = ["trace", "rt-tokio"] }
opentelemetry-otlp = { version = "0.15", features = ["trace", "http-proto", "reqwest-client", "tokio"] }
tokio = { version = "1", features = ["full"] } # For async runtime
uuid = { version = "1.0", features = ["v4"] } # For unique event IDs
futures-util = "0.3" # For WebSocket stream/sink
tokio-tungstenite = "0.21" # For WebSocket client
axum = "0.7" # For WebSocket server
tower-http = { version = "0.5", features = ["trace"] } # For Axum tracing middleware

```

#### src/main.rs (Illustrative, showing core setup and channel usage):

```

// src/main.rs
mod types; // Assuming your MyEvent and context helpers are in src/types.rs

use axum::{
    extract::{
        ws::{Message, WebSocket, WebSocketUpgrade},
        State,
    },
};

```

```

    response::Response,
    routing::get,
    Router,
};
use opentelemetry::{
    global,
    sdk::{
        trace::{self, Sampler},
        Resource,
    },
    KeyValue,
};
use opentelemetry_otlp::{self, WithExportConfig};
use tokio::sync::mpsc;
use tracing::{info, instrument, Span};
use tracing_subscriber::{prelude::*, registry::Registry, EnvFilter};
use types::{inject_current_context, MyEvent, extract_context_from_carrier}; // Import your
types

```

```

#[tokio::main]
async fn main() -> Result<(), Box<dyn std::error::Error>> {
    // 1. Configure OpenTelemetry Tracer Provider
    let otlp_exporter = opentelemetry_otlp::new_exporter()
        .http()
        .with_endpoint("http://localhost:4318/v1/traces") // OTLP Collector endpoint
        .with_timeout(std::time::Duration::from_secs(3))
        .build()?;

    let tracer = opentelemetry_otlp::new_pipeline()
        .tracing()
        .with_exporter(otlp_exporter)
        .with_trace_config(
            trace::config()
                .with_sampler(Sampler::AlwaysOn)
                .with_resource(Resource::new(vec![
                    KeyValue::new("service.name", "hybrid-rust-app"),
                    KeyValue::new("service.version", "0.1.0"),
                ])),
        )
        .install_batch(opentelemetry::runtime::Tokio)?;

    // 2. Create tracing-opentelemetry layer
    let opentelemetry_layer = tracing_opentelemetry::OpenTelemetryLayer::new(tracer);
}

```

```

// 3. Create tracing subscriber for console output and OpenTelemetry export
let subscriber = Registry::default()
    .with(EnvFilter::from_default_env())
    .with(opentelemetry_layer)
    .with(tracing_subscriber::fmt::layer().compact());

global::set_subscriber(subscriber)?;

info!("Application starting...");

// Setup channels for communication between async and non-async parts
let (ws_to_processor_tx, ws_to_processor_rx) = mpsc::unbounded_channel::<MyEvent>();
let (processor_to_ws_client_tx, processor_to_ws_client_rx) =
mpsc::unbounded_channel::<MyEvent>();

// Spawn non-async worker threads
for i in 0..2 { // Example: 2 worker threads
    let rx_clone = ws_to_processor_rx.clone(); // Clone receiver for each thread
    let next_tx_clone = processor_to_ws_client_tx.clone(); // Clone sender for each thread
    std::thread::spawn(move || {
        non_async_event_processor(rx_clone, next_tx_clone, i);
    });
}
// Drop the original receiver to ensure channels close when all senders are dropped
drop(ws_to_processor_rx);
drop(processor_to_ws_client_tx);

// Spawn async WebSocket client task
tokio::spawn(async move {
    // In a real app, you'd connect to an external WebSocket service here
    let (ws_stream, _) = tokio_tungstenite::connect_async("ws://echo.websocket.org")
        .await
        .expect("Failed to connect to external WebSocket");
    info!("Connected to external WebSocket service.");
    external_websocket_sender(processor_to_ws_client_rx, ws_stream).await;
});

// Setup Axum server
let app = Router::new()
    .route("/ws", get(websocket_handler))
    .with_state(ws_to_processor_tx) // Pass sender to Axum handlers

```

```

        .layer(tower_http::trace::TraceLayer::new_for_http()); // Basic HTTP tracing for Axum

let listener = tokio::net::TcpListener::bind("127.0.0.1:3000").await?;
info!("Axum server listening on http://127.0.0.1:3000");
axum::serve(listener, app).await?;

info!("Application finished.");
global::shutdown_tracer_provider(); // Ensure all traces are flushed
Ok(())
}

// --- Async Axum WebSocket Server ---
#[instrument(skip(ws, tx))]
async fn websocket_handler(
    ws: WebSocketUpgrade,
    State(tx): State<mpsc::UnboundedSender<MyEvent>>,
) -> Response {
    ws.on_upgrade(move |socket| handle_socket(socket, tx))
}

#[instrument(skip(socket, tx))]
async fn handle_socket(mut socket: WebSocket, tx: mpsc::UnboundedSender<MyEvent>) {
    info!("WebSocket connection established.");

    while let Some(msg) = socket.recv().await {
        if let Ok(msg) = msg {
            if let Message::Text(text) = msg {
                info!("Received message: {}", text);

                // --- INJECT CONTEXT: Async to Non-Async ---
                // Get the current OpenTelemetry context from the active tracing span (created by
                #[instrument] or TraceLayer)
                let trace_context_carrier = inject_current_context();

                let event = MyEvent {
                    id: uuid::Uuid::new_v4().to_string(),
                    payload: text,
                    trace_context_carrier,
                };

                if let Err(e) = tx.send(event) {
                    tracing::error!("Failed to send event to processing queue: {:?}", e);
                    break;
                }
            }
        }
    }
}

```

```

    }
  }
} else {
    tracing::warn!("WebSocket receive error or close.");
    break;
}
}
info!("WebSocket connection closed.");
}

// --- Non-Async Event Processor (Runs on dedicated threads) ---
fn non_async_event_processor(
    mut rx: mpsc::UnboundedReceiver<MyEvent>,
    next_tx: mpsc::UnboundedSender<MyEvent>,
    thread_id: usize,
) {
    info!("Non-async processor thread {} started.", thread_id);
    // Use `blocking_recv` for non-async context
    while let Some(event) = rx.blocking_recv() {
        // --- EXTRACT & ATTACH CONTEXT: Channel Receive ---
        let parent_otel_context = extract_context_from_carrier(&event.trace_context_carrier);

        // Make this context the current active context for the current thread.
        // All new tracing spans created within this block will be children of this context.
        let _guard = parent_otel_context.attach();

        // Create a new tracing span for the processing logic.
        // This span will automatically link to the parent context.
        let processing_span = tracing::span!(
            tracing::Level::INFO,
            "event_processing",
            event.id = &event.id,
            thread_id = thread_id,
            payload_len = event.payload.len()
        );
        let _processing_guard = processing_span.enter();

        info!("Thread {}: Processing event ID: {}", thread_id, event.id);
        // Simulate some CPU-bound work
        std::thread::sleep(std::time::Duration::from_millis(50));

        // Example: If this processing emits another event to another queue/thread:
        let new_payload = format!("Processed by Thread {}: {}", thread_id, event.payload);
    }
}

```

```

// --- INJECT CONTEXT: Non-Async to Async (or another Non-Async) ---
// Get the current OpenTelemetry context from the active span (processing_span)
let trace_context_carrier_for_next_event = inject_current_context();

let next_event = MyEvent {
    id: uuid::Uuid::new_v4().to_string(),
    payload: new_payload,
    trace_context_carrier: trace_context_carrier_for_next_event,
};

if let Err(e) = next_tx.send(next_event) {
    tracing::error!("Thread {}: Failed to send next event: {:?}", thread_id, e);
}

info!("Thread {}: Finished processing event ID: {}", thread_id, event.id);
// `_processing_guard` and `_guard` are dropped here, ending spans and detaching
context
}
info!("Non-async processor thread {} stopped.", thread_id);
}

// --- Async WebSocket Client (Sending to External Service) ---
#[instrument(skip(rx, ws_stream))]
async fn external_websocket_sender(
    mut rx: mpsc::UnboundedReceiver<MyEvent>,
    mut ws_stream:
tokio_tungstenite::WebSocketStream<tokio_tungstenite::MaybeTlsStream<tokio::net::TcpStrea
m>>,
) {
    info!("External WebSocket sender task started.");
    while let Some(event_to_send) = rx.recv().await {
        // --- EXTRACT & ATTACH CONTEXT: Channel Receive ---
        let parent_otel_context =
extract_context_from_carrier(&event_to_send.trace_context_carrier);
        let _guard = parent_otel_context.attach();

        // Create a span for the actual sending over WebSocket
        let send_ws_span = tracing::span!(
            tracing::Level::INFO,
            "external_websocket_send",
            event.id = &event_to_send.id,
            final_payload_len = event_to_send.payload.len()

```



```

);
let _send_ws_guard = send_ws_span.enter();

info!("Sending via external WebSocket: {}", event_to_send.payload);
let message = tokio_tungstenite::tungstenite::Message::Text(event_to_send.payload);
if let Err(e) = ws_stream.send(message).await {
    tracing::error!("Failed to send message over external WebSocket: {:?}", e);
}
info!("Message sent over external WebSocket.");
// `_send_ws_guard` and `_guard` are dropped here
}
info!("External WebSocket sender task stopped.");
}

```

### Step 3.3: Configure and Run OpenTelemetry Collector

This remains the same as in the previous explanation. The collector will receive OTLP traces and forward them to your Elastic APM Server.

#### **otel-collector-config.yaml:**

# otel-collector-config.yaml

receivers:

otlp:

protocols:

grpc:

http:

endpoint: 0.0.0.0:4318 # Default OTLP HTTP port

processors:

batch:

send\_batch\_size: 100

timeout: 1s

resource:

attributes:

- key: host.name

value: my-hybrid-app-host

action: upsert

exporters:

otlphttp/elastic:

endpoint: "http://<your-elastic-apm-server-host>:8200" # \*\*IMPORTANT: Replace with your

Elastic APM Server URL\*\*

compression: gzip

# headers:

```
# authorization: "Bearer <your_elastic_apm_api_key>"
```

service:

pipelines:

traces:

receivers: [otlp]

processors: [batch, resource]

exporters: [otlphttp/elastic]

### Run the collector:

```
./opentelemetry-collector-contrib --config otel-collector-config.yaml
```

## Step 3.4: Run Your Rust Application

cargo run

You can test this by connecting to `ws://127.0.0.1:3000/ws` using a WebSocket client (e.g., a browser's developer console or a tool like Postman/Insomnia) and sending text messages. You should see the traces appear in Kibana.

## 4. Explanation of Context Flow

### 1. Axum WebSocket Handler (`handle_socket`):

- An initial tracing span is created by `#[instrument]` or Axum's `TraceLayer`.
- When a message is received and an `MyEvent` is created: `inject_current_context()` is called. This function gets the `opentelemetry::Context` from the *current active span* (the Axum handler span) and serializes it into the `trace_context_carrier` `HashMap`.
- The `MyEvent` with the embedded context is sent to the `ws_to_processor_tx` channel.

### 2. Non-Async Event Processor (`non_async_event_processor`):

- When `rx.blocking_recv()` receives an `MyEvent`:
  - `extract_context_from_carrier()` deserializes the `HashMap` back into an `opentelemetry::Context`. This is the *parent context* from the Axum handler.
  - `parent_otel_context.attach()` makes this parent context the *current active context* for the worker thread.
  - A new `tracing::span!` (`event_processing`) is created. Because a parent context is active, this new span automatically becomes a *child* of the Axum handler span, linking them together in the trace.
- If this processor then emits a `next_event`:
  - `inject_current_context()` is called again. This time, it captures the context of the `event_processing` span (which is a child of the original Axum span).
  - This new context is embedded in the `next_event` and sent to the

processor\_to\_ws\_client\_tx channel.

### 3. **Async WebSocket Client (external\_websocket\_sender):**

- When rx.recv().await receives an MyEvent:
  - Similar to the non-async processor, extract\_context\_from\_carrier() extracts the context (which now represents the event\_processing span).
  - parent\_otel\_context.attach() makes it active.
  - A new tracing::span! (external\_websocket\_send) is created, automatically becoming a child of the event\_processing span.
- The actual WebSocket send operation happens within this span, completing the trace path from the initial incoming message to the final outgoing message.

By explicitly passing and activating the opentelemetry::Context at each boundary, you ensure that all segments of your event-driven flow are correctly linked together into a single, comprehensive distributed trace in Elastic.